Day 08 of 60

Stop watching tutorials.

Build THIS today.

The ultimate walkthrough to build a clean, end-to-end Machine Learning pipeline on the classic Titanic dataset. Learn how to write production-ready code that hiring managers actually value.

📖 6 slides · ⏱ 4 min read · 💻 Production code · 📌 Portfolio ready

1. Clean pipeline

No Jupyter notebook spaghetti. A modular, reproducible Python script.

2. Titanic solved

Best-practice preprocessing, scaling, modeling, and serialization.

3. Portfolio gold

The exact checklist of items to commit to GitHub to impress recruiters.

SLIDE 1 — THE TRAP

Jupyter Notebook Spaghetti

THE MISTAKE

Tutorial Copy-Pasting

Watching tutorials and copying cells into a single notebook gives you a false sense of achievement, but creates zero job-ready skills.

THE GOLD STANDARD

Modular pipelines

Hiring managers look for reproducible, self-contained Python scripts that automate data ingestion, cleaning, validation, and training.

Why this matters for your career

ML systems in production run as automated microservices, not Jupyter notebooks. Proving you can write structured pipeline scripts puts you ahead of 90% of applicants.

Ingestion & Preprocessing

Use Scikit-Learn's `Pipeline` and `ColumnTransformer` to clean and encode variables without data leakage.

pipeline.py

```python
import pandas as pd
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder

# 1. Load messy data & extract features
df = pd.read_csv("titanic.csv")
X = df[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare']]
y = df['Survived']

# 2. Define preprocessing sub-pipelines
num_pipeline = Pipeline([('imputer', SimpleImputer(strategy='median'))])
cat_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown='ignore'))
])

# 3. Combine transformers
preprocessor = ColumnTransformer([
    ('num', num_pipeline, ['Age', 'SibSp', 'Parch', 'Fare']),
    ('cat', cat_pipeline, ['Pclass', 'Sex'])
])
```

Why this code is production-grade

By packing steps into a `ColumnTransformer`, you ensure that numerical imputations (e.g. median age) and categorical encodings (e.g. sex mapping) are calculated during training and applied identically to new predictions without data leakage.

Train, Validate, & Export

```python
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
import joblib

# 4. Integrate preprocessing with modeling
pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('model', RandomForestClassifier(max_depth=5, random_state=42))
])

# 5. Split, train, and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
pipeline.fit(X_train, y_train)

# 6. Save model artifact for deployment
joblib.dump(pipeline, "titanic_model.joblib")
```

bash

```bash
$ python pipeline.py
[INFO] Ingesting raw titanic.csv... (891 passenger records loaded)
[INFO] Initializing features: numerical (Age, Fare, ...) & categorical (Sex, Pclass)
[INFO] Training RandomForestClassifier model on 712 samples...
✅ Model Validation Accuracy: 82.68%
[INFO] Exporting end-to-end model pipeline artifact to titanic_model.joblib
🎉 Pipeline complete! Saved model successfully in 0.38s
```

SLIDE 4 — GITHUB PROOF

Submit to Portfolio

Hiring managers look for a clean, modular repository, not a single notebook. Here is your portfolio checklist:

01 — PIPELINE.PY

Executable Script

A raw Python file that runs end-to-end with a terminal command and outputs validation metrics.

02 — SETUP & DEPS

Requirements.txt

Explicitly specified library versions (scikit-learn, pandas, joblib) so anyone can reproduce your pipeline.

03 — DOCUMENTATION

Readme.md

Explain the architecture, pipeline stages, evaluation scores, and steps to train and predict.

04 — INTERACTIVE UI

Streamlit App

Add a 15-line `app.py` Streamlit UI that loads your joblib file and makes predictions in real-time.

Recruiter impact

This shows software engineering discipline, attention to reproducibility, and deployment skills — a 10x stronger signal than a Kaggle score.

Follow Vinod Bavage for Day 09 →

SOON Tomorrow: **Preventing Data Leakage: The silent killer**

#MachineLearning #VinodBavage #MLPipeline

VINOD BAVAGE

VB **Vinod Bavage**
Machine Learning · AI

</> vinodbavage31

Day 08 / 60